

Micrometer

Spring Boot 3 Workshop - Devnexus 2023 - Atlanta

Jonatan Ivanov & Phil Webb

What is Observability?

- 3 pillars: Logging, Metrics, Distributed Tracing
- 4 pillars: + Events/Lineage(?)/Context/Metadata
- 6 pillars: + Profiles + Exceptions Arbitrary Wide Events, Signals
- But what about:
 - /health, /info, etc.
 - Service Registry/Discoverability
 - API Discoverability

What is Observability?



**How well we can understand the
internals of a system based on its
outputs**

(Providing meaningful information about what happens inside)

Why do we need Observability?

Today's systems are insanely complex (cloud)

(Death Star Architecture, Big Ball of Mud)



Why do we need Observability?

Environments can be chaotic

(You turn a knob here a little and services are going down there)

We need to deal with unknown unknowns

(We can't know everything)

Things can be perceived differently by observers

(Everything is broken for the users but seems ok to you)

Logging - Metrics - Distributed Tracing

Logging

What happened (why)?

Emitting events

Metrics

What is the context?

Aggregating data

Distributed Tracing

Why happened?

Recording causal ordering of events

Examples

Latency

Logging

Processing took 140ms

Metrics

P99.999: 140ms

Max: 150 ms

Distributed Tracing

DB was slow (lot of data was requested)

Examples

Error

Logging

Processing failed (stacktrace?)

Metrics

The error rate is 0.001/sec

2 errors in the last 30 minutes

Distributed Tracing

DB call failed (invalid input)

Checkpoint

Everyone knows that Observability \neq "Three Pillars" 😊

Logging with Spring

SLF4J + Logback

- SLF4J with Logback comes pre-configured
- SLF4J: Simple Logging Façade for Java
 - Simple API for logging libraries
- Logback Natively implements the SLF4J API
- If you want Log4j2 instead of Logback:
 - `spring-boot-starter-logging`
 - + `spring-boot-starter-log4j2`

Setup Logging

Add org property

- We will need something that we can use to query:
 - All of our apps (`spring.application.org`)
 - Only one app (`spring.application.name`)
 - Only one instance (we only have one instance per app)

spring:

application:

name: dog-service

org: petclinic

Setup Logging

Add Loki4J

- Copy
From: `dog-client/src/main/resources/logback-spring.xml`
To: `dog-service/src/main/resources/logback-spring.xml`
- Add dependency to `pom.xml`

```
<dependency>  
  <groupId>com.github.loki4j</groupId>  
  <artifactId>loki-logback-appender</artifactId>  
  <version>1.4.0</version>  
</dependency>
```


Setup Logging

Do we have logs?

- Got to Grafana: <http://localhost:3000>
- Choose Explore, then Loki from the drop down
- Search for `application = dog-service`
- Search for `org = petclinic`
- We will get back to our logs later

Checkpoint

Everyone has logs in Loki for both services

Metrics with Spring

Micrometer

- Popular Metrics library on the JVM
- Like SLF4J, but for metrics
- Simple API
- Supports the most popular metric backends
- Comes with **spring-boot-actuator**
- Spring projects are instrumented using Micrometer
- A lot of third-party libraries use Micrometer

Metrics with Spring

Like SLF4J, but for metrics ...

AppOptics

Ganglia

OpenTSDB

Atlas

Graphite

OTLP

Azure Monitor

Humio

Prometheus

CloudWatch (AWS)

InfluxDB

SignalFx

Datadog

JMX

Stackdriver (GCP)

Dynatrace

KairosDB

StatsD

Elastic

New Relic

Wavefront (VMware)

(/actuator/metrics)

Sidetrack: Observation API

You want to instrument your application...

- Add logs (application logs)
- Add metrics
 - Start/Stop Timers, Increment Counters
- Add Distributed Tracing
 - Start/Stop Spans
- Log Correlation
- Context Propagation

Observation API (Micrometer 1.10)

```
Observation observation = Observation.start("talk", registry);

try { // TODO: scope
    Thread.sleep(1000);
}
catch (Exception exception) {
    observation.error(exception);
    throw exception;
}
finally { // TODO: attach tags (key-value)
    observation.stop();
}
```

Observation API (Micrometer 1.10)

```
ObservationRegistry registry = ObservationRegistry.create();

registry.observationConfig()
    .observationHandler(new MeterHandler(...))
    .observationHandler(new TracingHandler(...))
    .observationHandler(new LoggingHandler(...))
    .observationHandler(new AuditEventHandler(...));

Observation observation = Observation.start("talk", registry)
    // let the fun begin...
observation.stop();
```

Observation API (Micrometer 1.10)

```
Observation.createNotStarted("talk", registry)
    .lowCardinalityKeyValue("conference", "DN")
    .highCardinalityKeyValue("uid", userId)
    .observe(this::talk);
```

@Observed

Setup Metrics

Add the org to Observations

`...actuator.ActuatorConfiguration.java`

```
@Bean
```

```
ObservationFilter orgFilter(  
    @Value("${spring.application.org}") String org) {  
    return context -> context.addLowCardinalityKeyValue(  
        KeyValue.of("org", org)  
    );  
}
```

Setup Metrics

Let's check Metrics

- Go to `http://localhost:8080/actuator/prometheus`
- 401 🤔
- So Prometheus is broken too? `http://localhost:9090/targets`
- Spring Security! 👁️
- Let's disable it, what could go wrong!? 😈
- Everything, please don't do this in prod!
Except you want everyone know about it. 😈

Setup Metrics

Disable auth for certain endpoints

`SecurityConfiguration.java`

`requests`

```
.requestMatchers("/dogs", "/actuator/**").permitAll();
```

Setup Metrics

Add org and application tags to every meter

- Why only a few Prometheus time series has the org tag?
- Not everything is created through the Observation API e.g.: heap utilization
- Let's add tags to everything!

management:

metrics:

tags:

application: \${spring.application.name}

org: \${spring.application.org}

Setup Metrics

Remove Observation name customization

- We are going to depend on default behavior
- So let's remove the custom http observation renaming
- Remove/comment out

```
# observations:
#     http:
#         server:
#             requests:
#                 name: "http.server.in"
```

Setup Metrics

Add histogram support for http metrics

- We want to see the latency distributions on our dashboards
- We want to calculate percentiles (P99?)

```
management:
```

```
  metrics:
```

```
    distribution:
```

```
      percentiles-histogram:
```

```
        # all: true
```

```
        http.server.requests: true
```

Setup Metrics

Let's check the HTTP and JVM metrics

- Let's check `/actuator/metrics`
`/actuator/metrics/{metricName}`
`/actuator/metrics/{metricName}?tag=key:value`
- Let's write a Prometheus query (`HELP.md`)

`sum by (application)`
`(rate(http_server_requests_seconds_count[5m]))`
- Let's check the dashboards: go to Grafana, then Browse
 - Spring Boot Statistics
 - Dogs

Checkpoint

Everyone has metrics on the dashboards

Distributed Tracing with Spring

Micrometer Tracing and Spring Cloud Sleuth

- **Boot 2.x** - Spring Cloud Sleuth
- **Boot 3.x** - Micrometer Tracing (Sleuth w/o Spring dependencies)
- Provide an abstraction layer on top of tracing libraries
 - Brave (OpenZipkin), default
 - OpenTelemetry (CNCF), experimental
- Instrumentation for Spring Projects, third-party libraries, your app
- Support for various backends

Setup Distributed Tracing

Add Micrometer Tracing Dependencies

```
<dependency>  
  <groupId>io.micrometer</groupId>  
  <artifactId>micrometer-tracing-bridge-brave</a  
</dependency>  
<dependency>  
  <groupId>io.zipkin.reporter2</groupId>  
  <artifactId>zipkin-reporter-brave</artifactId>  
</dependency>
```

Setup Distributed Tracing

Set sampling probability

```
management:
```

```
  tracing:
```

```
    sampling:
```

```
      probability: 1.0
```

Setup Distributed Tracing

Setup log correlation

Copy this from dog-client's `application.yml`

`logging:`

`pattern:`

`level: "%5p [%${spring.application.name:}, %X{traceId:-}, %X{spanId:-}]"`

`level: org.springframework.web.servlet.DispatcherServlet: DEBUG`

Setup Distributed Tracing

Let's look at correlated logs

```
[  
  dog-service,  
  641e50ff9c6911eb4d96f9d4a19ebc82,  
  c3e2ea2a9b4a3f60  
]
```

```
[app-name, traceId, spanId]
```

Setup Distributed Tracing

Let's look at some traces


- Go to Grafana, then Explore and choose Tempo
- Terminology
 - Span
 - Trace
 - Tags
 - Annotations

Checkpoint

Everyone has log correlation and traces in Tempo

Interoperability

Logs -> Trace / Trace -> Log



Setup Observations

Disable Spring Security Observations

`ActuatorConfiguration.java`

`@Bean`

```
ObservationPredicate noSpringSecurityObservations() {  
    return (name, ctx) -> !  
        name.startsWith("spring.security.");  
}
```

Setup Observations

Disable Actuator Observations

`ActuatorConfiguration.java`

```
if (name.equals("http.server.requests") &&  
    ctx instanceof ServerRequestObservationContext sc) {  
    return !sc.getCarrier()  
        .getRequestURI()  
        .startsWith("/actuator");  
}
```

Setup Observations

Enable JDBC Observations

- Tadaa Tsuyukubo 🕶️
- `net.ttddyy.observation:datasource-micrometer-spring-boot
(1.0.1)`

`jdbc:`

`datasource-proxy:`

`include-parameter-values: true`

`query:`

`enable-logging: true`

`log-level: INFO`

Setup Observations

Add custom Observation


OwnerService.java

Observation

```
.createNotStarted("getDogs", registry)
.contextualName("gettingOwnedDogs")
.highCardinalityKeyValue("owner", owner)
.observe(() -> {
    //...
});
```


Interoperability

Logs \Leftrightarrow Traces \Leftrightarrow Metrics



Interoperability

How to check Exemplars

- Exemplars are only available if you request the OpenMetrics format
- Your browser does not do this

```
http :8081/actuator/prometheus
```

```
'Accept: application/openmetrics-text;version=1.0.0'
```

```
(| grep trace_id)
```

Checkpoint

Logs <-> Metrics <-> Traces

Setup Observations

Log error and signal it

OwnerController.java

```
ProblemDetail onNoSuchDogOwner (
    HttpServletRequest request,
    NoSuchDogOwnerException ex) {

    logger.error("Ooops!", ex);
    ServerHttpObservationFilter
        .findObservationContext(request)
        .ifPresent(context -> context.setError(ex));
}
```

Setup Observations

Hack error reporting for Tempo

```
ObservationFilter tempoErrorFilter() {
    return context -> {
        if (context.getError() != null) {
            context.addHighCardinalityKeyValue(
                KeyValue.of("error", "true")
            );
            context.addHighCardinalityKeyValue(
                KeyValue.of(
                    "errorMessage", context.getError().getMessage()
                )
            );
        }
        return context;
    };
}
```

Setup Observations

Hack DB tags for Tempo ServiceGraph

```
if (ctx instanceof DataSourceBaseContext dsCtx) {  
    ctx.addHighCardinalityKeyValue(  
        KeyValue.of(  
            "db.name", dsCtx.getRemoteServiceName()  
        )  
    );  
}
```

Actuator!

Add Java and OS InfoContributors

management:

 info:

 java:

 enabled: true

 os:

 enabled: true

Checkpoint

Everything works! 🤙




James Ward



@_JamesWard



Upgraded another service to JDK 17 and holy crap these out-of-the-box improvements in GC are 🤯



9:29 AM · Mar 28, 2023 · 82.3K Views

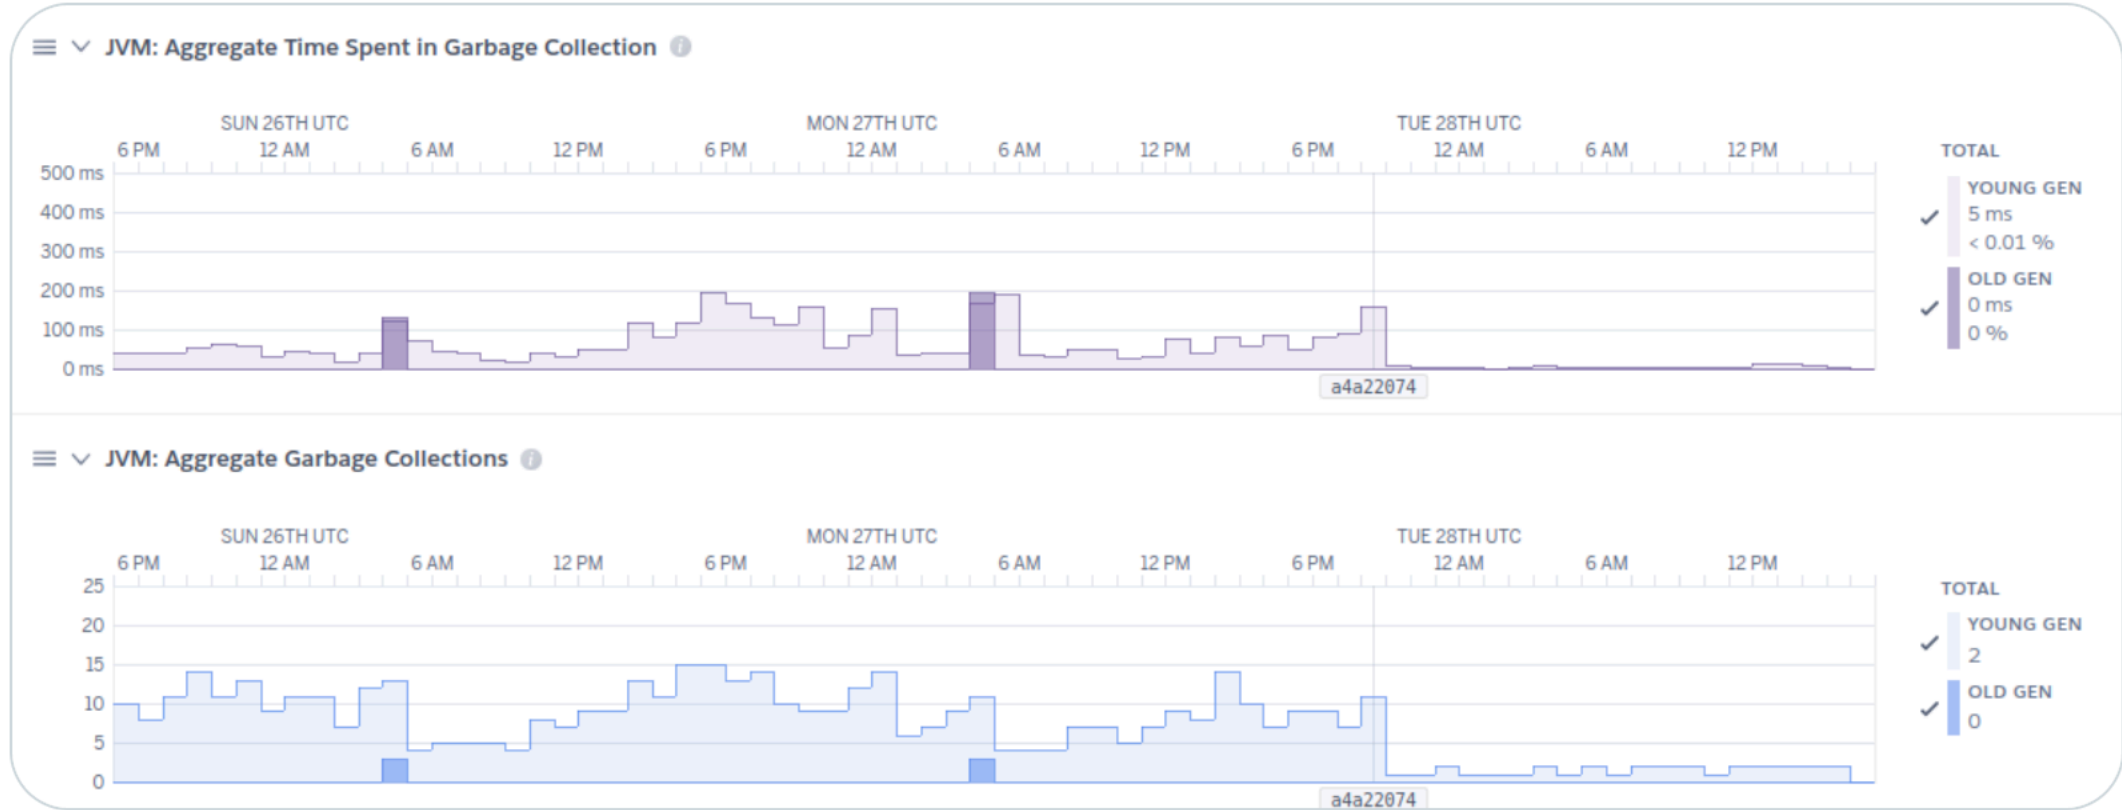
89 Retweets 7 Quotes 612 Likes 52 Bookmarks



James Ward ✓
@_JamesWard



Upgraded another service to JDK 17 and holy crap these out-of-the-box improvements in GC are 🤯



9:29 AM · Mar 28, 2023 · 82.3K Views

89 Retweets 7 Quotes 612 Likes 52 Bookmarks



Jonatan Ivanov @jonatan_ivanov · Mar 28
Replying to @_JamesWard and @spencerbibb
Is it G1-G1? Or what is the GC before/after?



Extra

Inject Latency

```
docker exec toxiproxy  
  /toxiproxy-cli toxic add  
  --toxicName base-latency  
  --type latency  
  --downstream  
  --toxicity 1.0  
  --attribute latency=50  
  --attribute jitter=0  
dog-db
```

Extra

- avg and p99
- The memory leak that only existed in prod
- The memory leak that was not a memory leak
- The zombie bot
- The hidden service from the past
- Auth failure (clock skew)
- Trade offs of Java Agents (for instrumentation)
- High Cardinality - Friend or Foe?
- How Not to Measure Elapsed Time

Extra

JVM ergonomics vs. Containers

- cgroups + namespaces
- Memory rq/limit: heap size
- CPU rq
 - Number of GC threads and GC algorithm
 - Number of runtime compiler threads (JIT)
 - Common Pool size (ForkJoinPool, Parallel Streams)
 - 3rd party thread Pools: `Runtime#availableProcessors` (Little's law)

Extra TeaHouse

- `actuator/health` and `actuator/info` (TeaHouse)
- Service Registry/Discoverability (Eureka, Spring Boot Admin)
- API Discoverability (Swagger, HATEOAS)
- Access logs, Logbook, GC logs
- FlyWay

Extra

actuator/health and actuator/info (TeaHouse)

- Wrong version deployed
- Right version is deployed but it was not build against the commit you thought
- Wrong environment
- Wrong timezone or local
- Wrong certificate/cert-chain
- Unpatched OS
- The memory leak that only existed in prod

Checkpoint

Observability done!